Firing Alerts

Overview

API Reachability: UP

Model Status: LOADED

MongoDB Reachability: UP

API Uptime: 16.2 mins

Predict Throughput: 0.0231 req/s

Firing Alerts: 4

Inference

Inference Latency

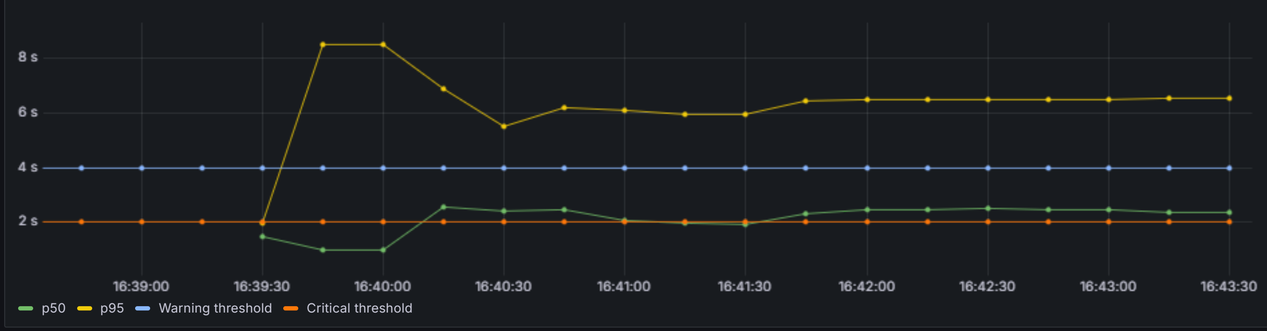

Latency p95: 6.79 s

Median Confidence: 94.1%

High-Risk Rate: 0.0873 req/s

Confidence Trend

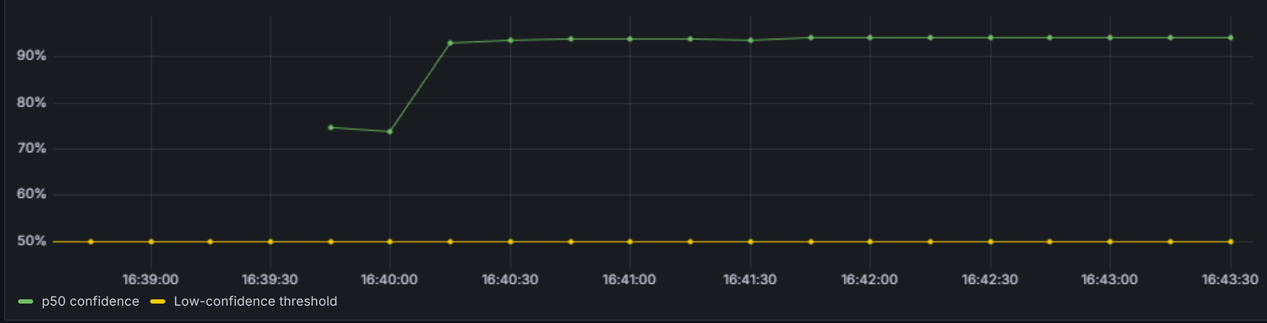

High-Risk Prediction Spike

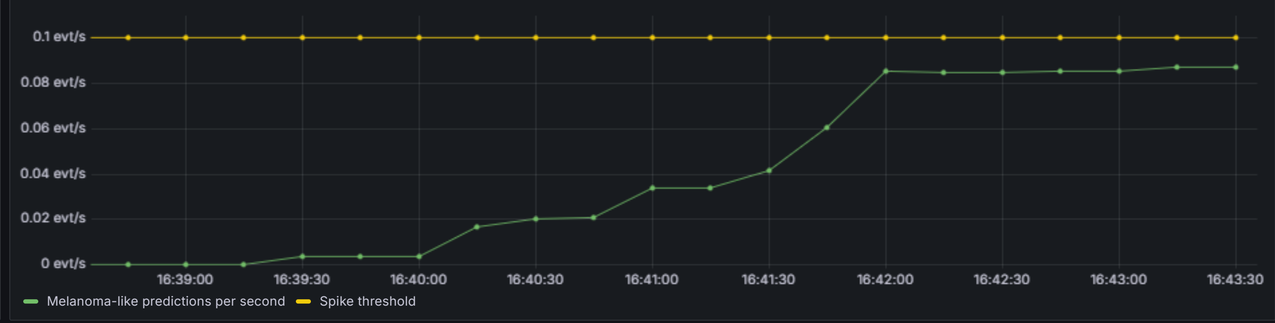

API And Traffic

API Error Ratio: 0%

Predict Rate: 1.39 req/m

Negative Feedback Ratio: 0

Total Request Rate: 0.0231 req/s

Request Rate By Endpoint

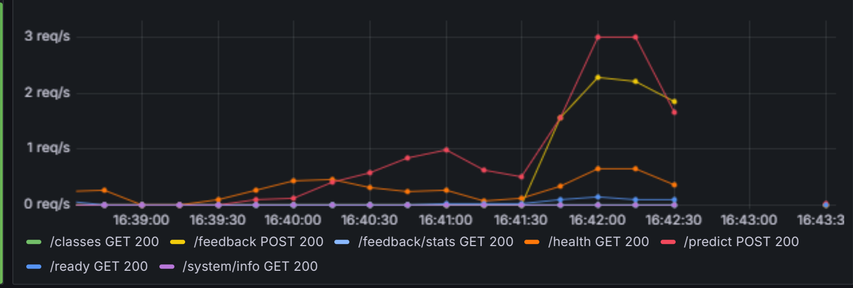

Error Ratio Trend

Negative Feedback Spike

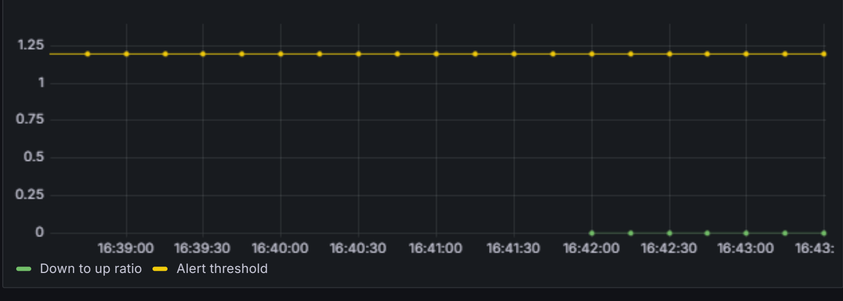

Traffic Alert Thresholds

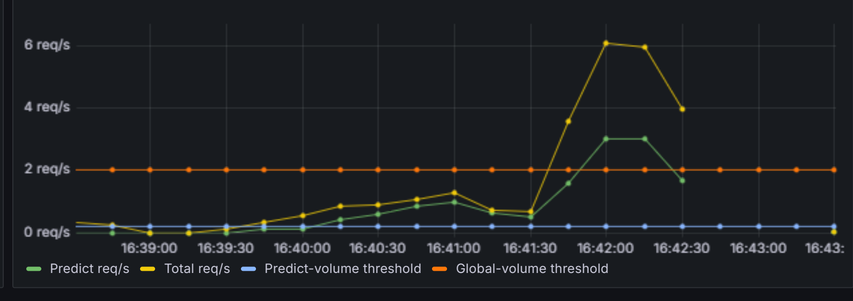

Resources

System CPU: 83.5%

System Memory: 80.8%

Request Memory p95: 650 MB

Active Requests: 0

GPU Allocated: 0 MB

GPU Reserved: 0 MB

System CPU Trend

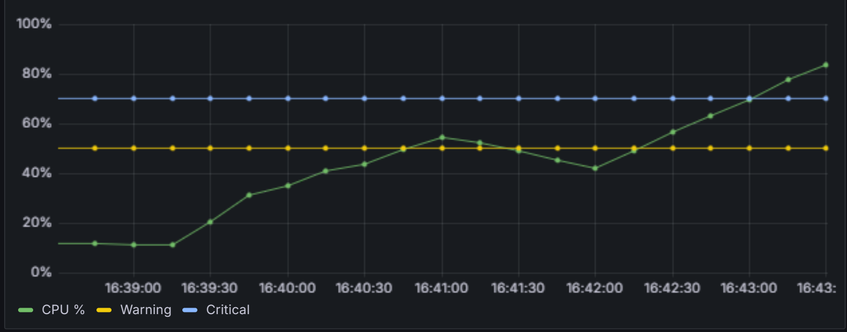

System Memory Trend

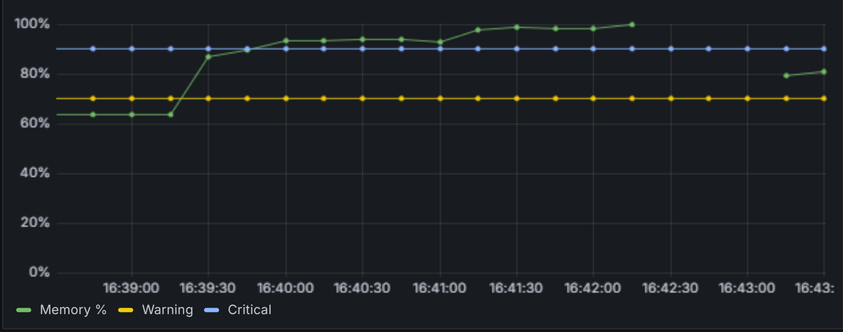

Request Memory Trend

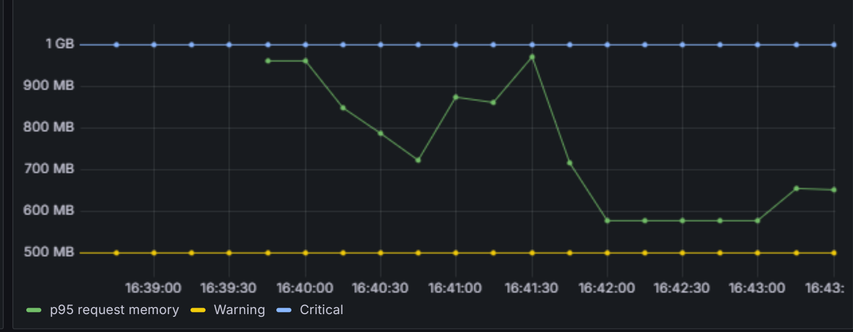

Database And Abuse

MongoDB Availability

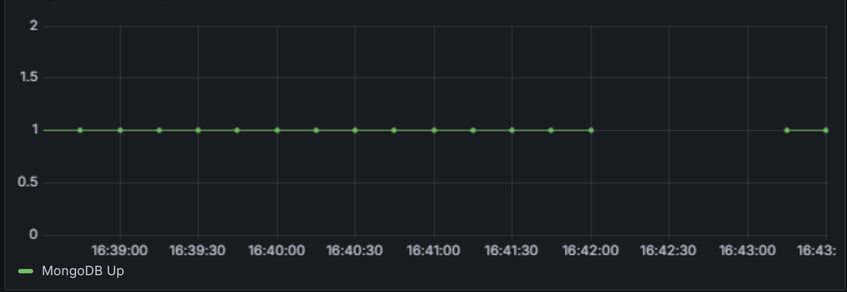

Registrations In 5m: 49.4

User Rate Limit Breaches

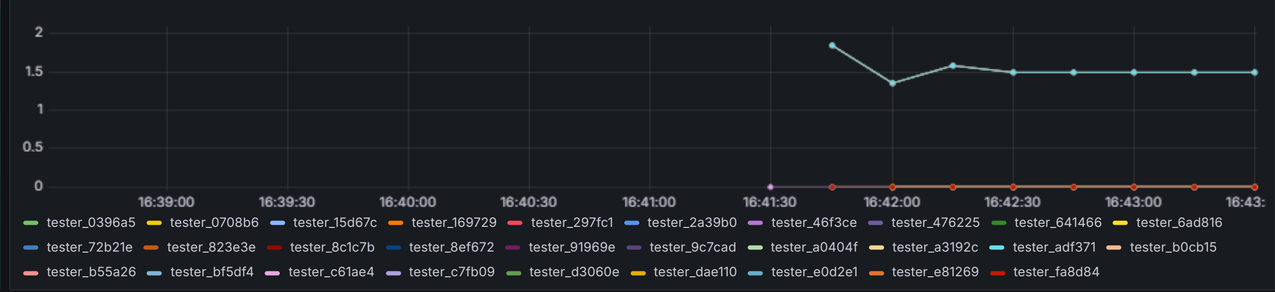

Alerts

Firing Alerts By Rule

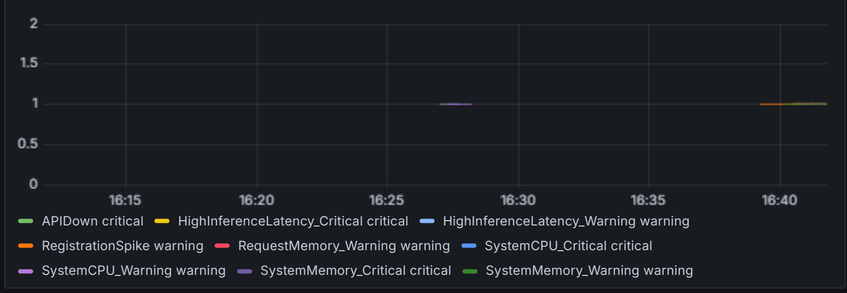

Firing Alerts By Severity

Active Alerts

- HighInferenceLatency_Critical — View alert rule
  Firing for 3m 56s
  › 1 instance
- HighInferenceLatency_Warning — View alert rule
  Firing for 3m 56s
  › 1 instance
- RegistrationSpike — View alert rule
  Firing for 5m 33s
  › 1 instance
- RequestMemory_Warning — View alert rule